

Monitoring Hadoop: Practical exercise

You will be working in pairs and using a preconfigured Linux VM (virtual machine) with Hadoop (v1.20), PIG and HIVE already installed.

Python and Java are available on the Linux box.

A number of text editors are available, including nano, vi and vim.

You will be working completely at the command line so you will need to access the VM via SSH using an SSH client application (such as PuTTY if you are using Windows).

The VMs are at:

`shu-hadoopn.cloudapp.net`

There are two preconfigured users: `hduser` and `azureuser`

To perform the Hadoop exercises, you will need to login as `hduser`, so:

```
ssh hduser@shu-hadoopn.cloudapp.net
```

(Don't forget to replace `n` with the number for your group).

Hadoop Web UIs are also available (network firewalls permitting):

`http://shu-hadoopn.cloudapp.net:50070/` – web UI of the NameNode daemon `http://shu-hadoopn.cloudapp.net:50030/` – web UI of the JobTracker daemon `http://shu-hadoopn.cloudapp.net:50060/` – web UI of the TaskTracker daemon

Part 1: Getting ready

1. Some sample data has been made available to you in the `gutenberg` directory (this is located at `/home/hduser/gutenberg`). It has already been loaded into the VM's HDFS.
2. Write a simple Java `mapper` and `reducer` to do a wordcount on the contents of these three files.
3. Use the Linux `time` command to time how long the wordcount takes. See the `guidance notes` below for an example of how to do this.
4. Repeat exercise 2 but write the `mapper` and `reducer` in Python and use `Hadoop streaming` to run the wordcount.
5. Again, use the `time` command to time how long this takes.

Guidance Notes

- You will find some sample Python scripts to use in the `python-streaming` directory. This is located at `/home/hduser/python-streaming`
- If you are unfamiliar with how to run Hadoop jobs from the command line, follow Michael Noll's tutorials:
- Running a MapReduce job:
`http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/#running-a-mapreduce-job`
- Running a Hadoop streaming Python job:
`http://www.michael-noll.com/tutorials/writing-an-hadoop-mapreduce-program-in-python/#running-the-python-code-on-hadoop`
- The Linux `time` command: `http://unix.stackexchange.com/questions/10745/how-do-i-time-a-specific-command`
- Basic usage is: `time <command>` and you will get something back like (it's the `real` time you are interested in:

```
real    0m0.178s
user    0m0.003s
sys     0m0.005s
```

Part 2: Comparing Native Java, Python streaming, PIG and HIVE

For these exercises, you have been given some data about airline departures from airports in the USA for 2008. The data is in the `airline_data` directory (`/home/hduser/airline_data`) and is described in the `README` which can be read using a text editor such as `nano`.

1. Put the data into a suitable directory in the HDFS. Instructions on how to do this are in the `Michael Noll` tutorial mentioned in the previous section.
2. Write a native Java `mapper` and `reducer` to calculate the average delay time across all flights.
3. Run this job and time how long it takes using the Linux `time` command.
4. Write and execute a Python streaming `mapper` and `reducer` to calculate the average delay. Time how long it takes.
5. Write a PIG script to calculate the average delay. Use your preferred method to run and time the script processing time (you can do this either interactively or by running the script).
6. Write and execute a HIVE query to calculate the average delay. Use your preferred method to run and time the script processing time.

Guidance Notes

Some Web resources that might help:

Java MapReduce calculate averages: Example code to demonstrate an approach: `https://gist.github.com/rishav-rohit/8133433`

Python MapReduce calculate averages: This one is for you to work out!

PIG calculate averages: `http://stackoverflow.com/questions/15212985/calculate-average-using-pig`

HIVE calculate averages:

`http://stackoverflow.com/questions/23804281/hive-grouping-and-calculating-average-by-calculating-distinct`

Part 3: Finishing off

Some points to consider to feedback at the end of the session:

1. What factors can be changed to improve the execution time of a job?
2. How can you determine which factors to change?
3. What is meant by `total time to solution` ? How can this concept help determine the most appropriate solution?